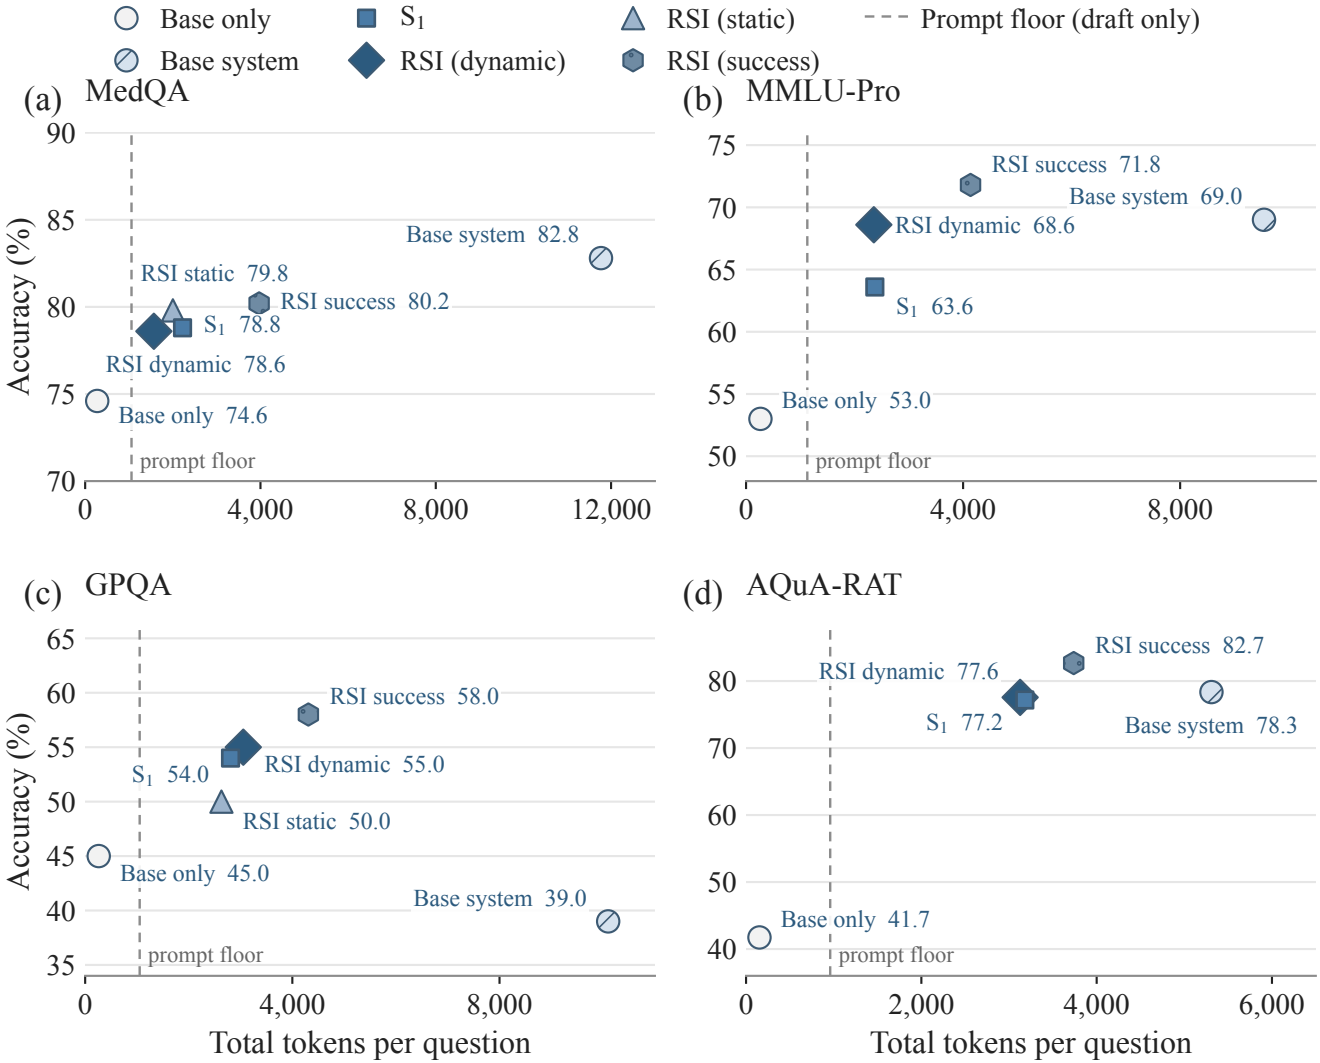

Locked tests (MedQA 500, MMLU-Pro 500 test, GPQA 100, AQuA 254); tokens = manager + advisor prefill and output per question. RSI: v3 + GRPO, round 3. Base only: base model answers with the letter only (16-token budget). S_1 tokens on MMLU-Pro and AQuA from the v3 S_1 run. RSI (success) tokens from the v1 success runs (same labelling rule) except GPQA, where the RSI run itself was measured. RSI (static) omitted where its tokens were not measured.