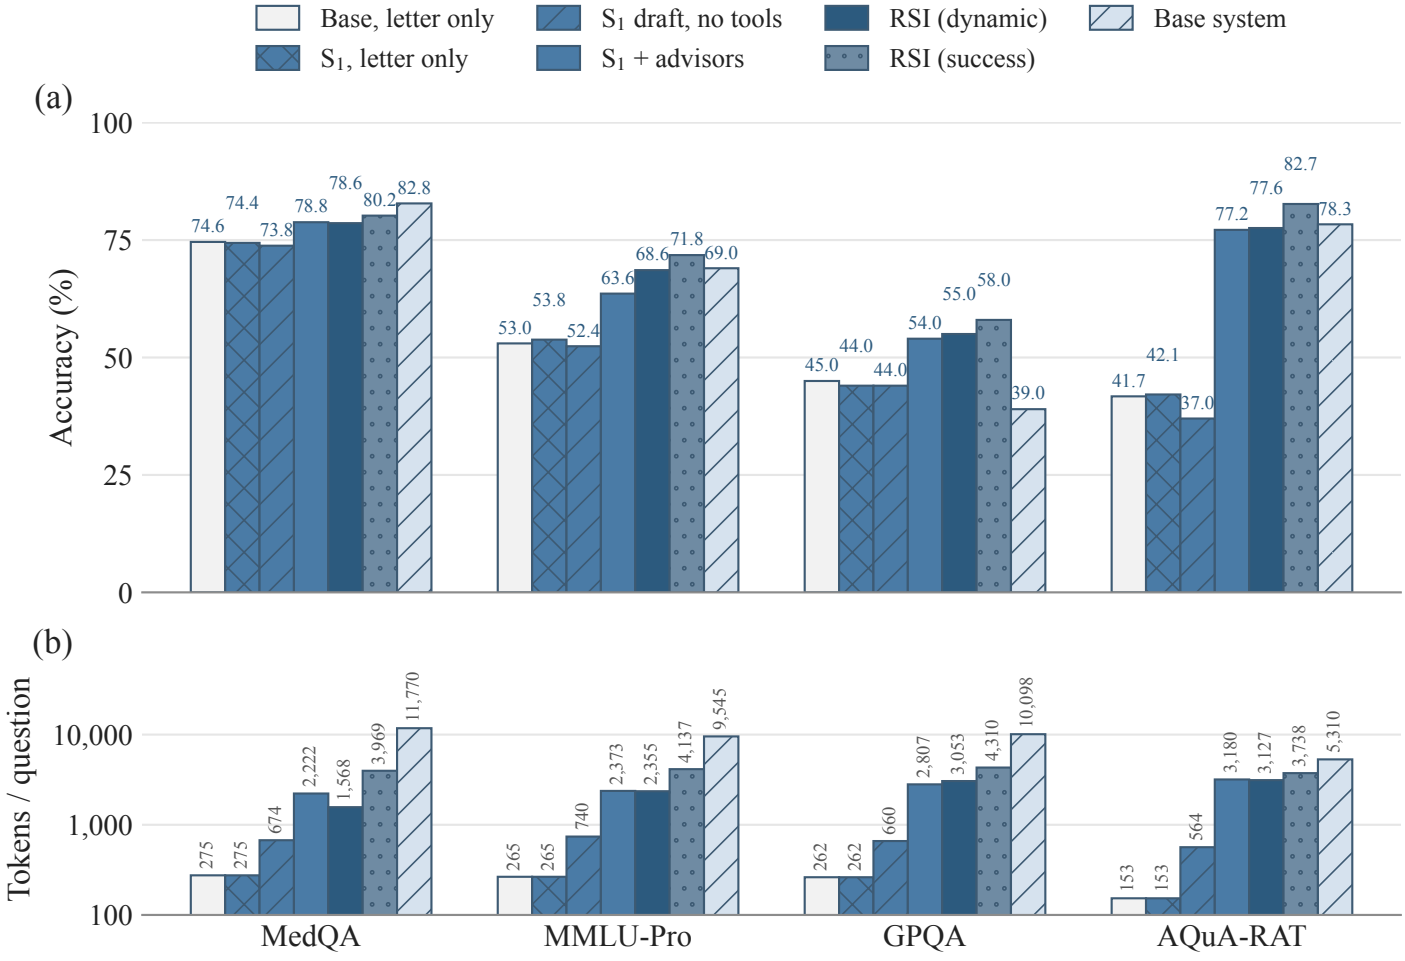

Locked tests: MedQA 500, MMLU-Pro 500 (test), GPQA 100, AQuA-RAT 254. (a) accuracy; (b) total tokens per question, log axis. Letter-only = model must answer with the letter (16-token budget, forced choice, 100% parseable); base system = untrained manager with tools. Tokens = prefill + output; S_1 + advisors on MMLU-Pro and AQuA-RAT from the v3-chain run of the same S_1; RSI (success) on MedQA, MMLU-Pro and AQuA-RAT from the v1 success runs (same labelling rule, call rate 0.93-1.00).